

Summary:

Backend engineer with 4 years of experience in automation and backend systems, building scalable asynchronous pipelines and production-grade APIs.

Skills:

Primary: Python, FastAPI, Flask, Kafka, AsyncIO, REST API Design, PostgreSQL, Redis

Secondary: JavaScript, Docker, AWS (ECS, RDS, MSK, ElastiCache, S3), GCP (Cloud Run/Build), Terraform, Git, Linux

Experience:

AVRL.io [via Husys Consulting Limited] – Software Engineer

Oct 2021 – September 2025

- Designed and maintained RPA-style automation systems using JavaScript and Python to automate high-frequency logistics and freight bidding workflows, significantly reducing manual operational effort.
 - Built and deployed automation scripts across client-managed Windows VMs, ensuring stability, fault tolerance, and production reliability in heterogeneous environments.
 - Containerized a custom browser-based automation stack using Docker and prototyped a Python Playwright-based replacement, improving portability, maintainability, and testability of automation workflows.
 - Contributed to a Flask-based internal onboarding platform (REST APIs, database design), gaining hands-on experience in backend API development and persistence layers.
-

Selected Personal Projects:

Real-time Feature Store Pipeline — TSSPDCL Accident Data

[\[GitHub\]](#)

- Scraped 10,000+ electrical accident records from the TSSPDCL government portal (2006–2026) with a rate-limited, idempotent scraper; migrated to RDS Postgres as the source of truth.
 - Built a Kafka pipeline (Amazon MSK) with a Python producer streaming batches via a Redis-tracked cursor and an AsyncIO consumer computing ~750 aggregated features into an ElastiCache Redis feature store.
 - Designed the feature schema across 6 dimensions — temporal, geographic (15 divisions), accident type, voltage class, and compensation pipeline state — served via FastAPI on ECS Fargate.
 - Integrated an Anthropic LLM inference endpoint that receives full feature context on each request, enabling natural language queries over the computed feature store.
 - Provisioned full AWS infrastructure (VPC, RDS, MSK Serverless, ElastiCache, ECS Fargate, ALB) using Terraform; shipped a live dashboard with replay controls, real-time feature visualizations, and an embedded AI chat interface.
-

Secure SFTP Ingestion & Processing Pipeline

[\[GitHub\]](#)

- Designed and implemented a secure, multi-vendor batch ingestion pipeline using SFTP (AWS Transfer Family) and PGP encryption.
 - Built a producer service to generate, encrypt, and atomically upload vendor files using GPG and SFTP with SSH key authentication.
 - Developed a consumer service that polls SFTP, decrypts files, and guarantees idempotent processing via SQLite-backed state tracking and retries. Integrated asynchronous AI-based document interpretation using Redis Queue (RQ), producing structured JSON outputs for downstream consumption.
 - Exposed a Flask-based monitoring API with health checks, file state inspection, filtering, and processing statistics.
-

Certifications:

- **MLOps | Machine Learning Operations Specialization** – Coursera, Duke University
 - **AWS** Developer Associate, **GCP** Associate Cloud Engineer
-

Education:

Geethanjali College of Engineering & Technology, B. Tech

Hyderabad, TS